

Part 5: Deadlocks and Starvation

- **Deadlock Problem**
- System Model
- Deadlock Conditions
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection – **not examinable**

Operating Systems 5.1 Part 5 Deadlock and Starvation

The Deadlock Problem

- **Deadlock:** A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set
- **Example 1:**
 - System has two different types of tape drives; **both P1 and P2 need the two tapes to finish execution**
 - P_1 and P_2 each hold one tape drive and need the other
- **Example 2:** Semaphores A and B , initialized to 1


```

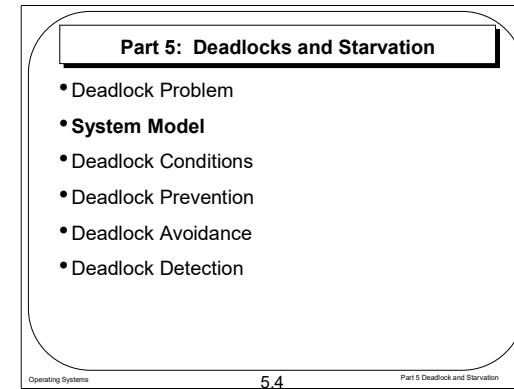
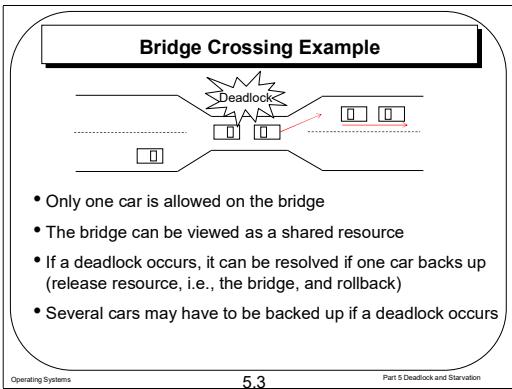
 $P_0$             $P_1$ 
wait (A);       wait (B);
wait (B);       wait (A );
      
```

Context switch points

Operating Systems 5.2 Part 5 Deadlock and Starvation

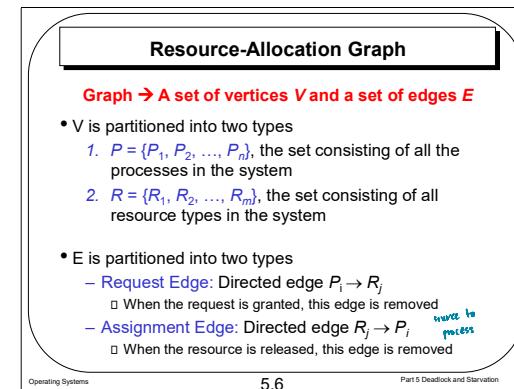
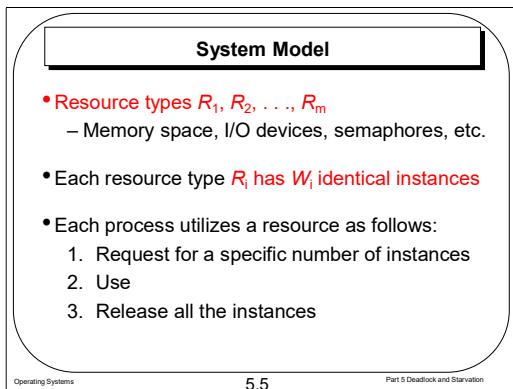
We have already seen deadlocks in the previous chapter. A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.

Another example is incorrect usage of semaphores. Say, we have two semaphores A and B initialized to 1. The two processes P_0 and P_1 can deadlock. We have seen this example in the previous chapter.



The bridge is a resource and the cars are processes.

Deadlock occurs if there are two cars on the bridge. When deadlock occurs, it can be resolved if one car backs up. Note, several cars may have to be backed up if a deadlock occurs. That means, deadlock recovery may affect other processes. Not only the process in the deadlock state are affected, but also other processes that have dependency on the deadlocked process.



We now introduce the system model for deadlock. In this model, we have different resource types $R_1 \dots R_m$. The resource type can be memory space or I/O devices. Each resource type R_i has W_i identical instances. For example, we have two hard disks in the system. R_i is the hard disk and $W_i=2$. the instances are all identical.

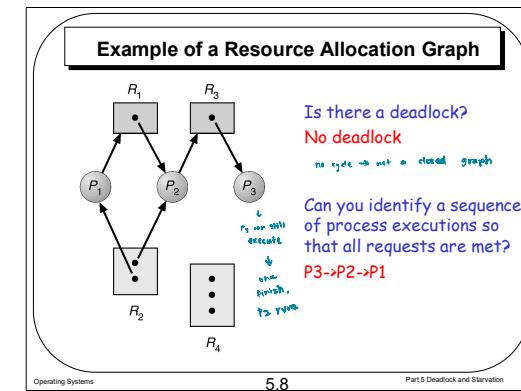
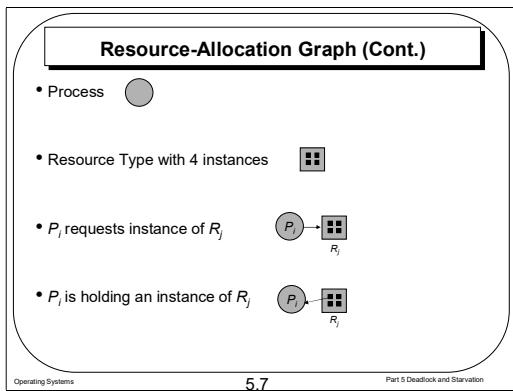
A process utilizes a resource as follows: first sends a request of a certain number of instances.

If the instances are available, then the request is granted. Otherwise, the process needs to wait. If the request is granted, the process uses the resource, and after some time, releases the resource.

Consider the system has multiple processes and multiple resource types. How to model resource request, use and release?

We use a data structure called the resource-allocation graph to model the resource allocations in the system. The graph has a set of vertices V and a set of edges E . V is partitioned into two types: process vertex and resource type vertex. The process vertices represent all the processes, and resource vertices represent all the resource types in the system.

When a process sends a request for R_j , we add a directed edge from P_i to R_j . *When the request is granted, the request edge is removed. Then, we add an assignment edge between the two vertices in the reverse direction. When the resource is released, the assignment edge is removed.*



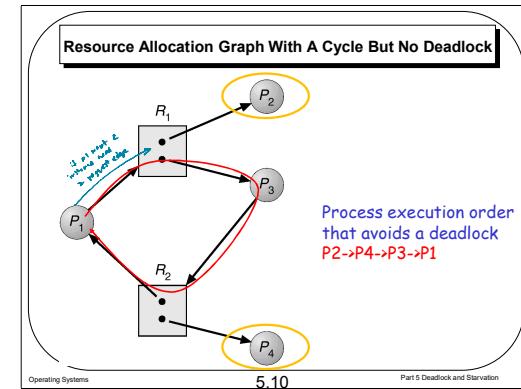
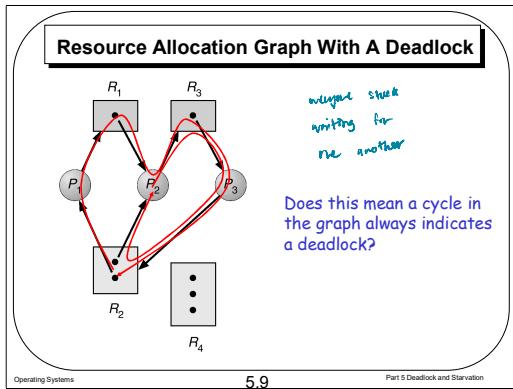
The number of solid black dots within the rectangle represents the number of instances for the resource.

When an instance of R_j is assigned to P_i , we draw an edge from one dot within the rectangle to P_i . Assignment of an instance of resources to the process: a certain dot within the resource vertex to the process vertex.

We have four resource types. Three processes.

We can actually examine the structure of the resource allocation graph to see whether the system has deadlock or not. Ideally, if the system has deadlock, the resource allocation must have a cycle. why? By the definition of deadlock ☺

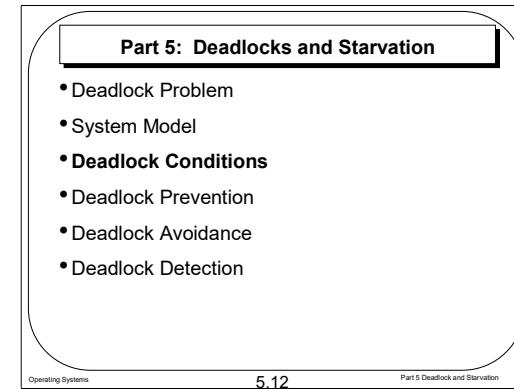
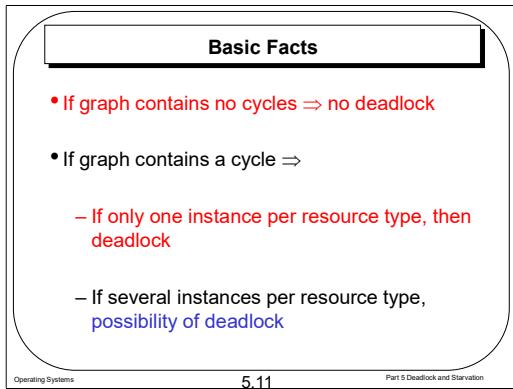
In this case there is no deadlock. Can you find a sequence of process execution that all the processes finish execution? If P_3 finishes, P_2 can finish, and then P_1 can finish.



How many cycles are there in this resource allocation graph?

Two cycles. First cycle: P2 waits for P3 and P3 waits for P2. Second cycle: P1 waits for P2, P2 waits for P3, and P3 waits for both P1 and P2. In this case, no process can make any progress. →Deadlock.

One cycle. But no deadlock. Because P2 can finish and release R1. Then P4 can finish. Then, P3 and P1 can finish.



So, in summary, if a graph contains no cycle, no deadlock for sure. Because we do not have the problem of A waiting for B and B waiting for A.

If the graph contains a cycle, we need to look at the resource types. If only one instance per resource type, then there is a deadlock. Because no process can make any progress. If several instances per resource type, deadlock is possible, but not 100% sure.

We will introduce banker's algorithm to solve the problem with multiple instances later on.

Deadlock Conditions

Deadlock **may** arise if the following four conditions hold **simultaneously** (necessary, but not sufficient)

1. **Mutual exclusion:** Only one process at a time can use a resource instance
2. **Hold and wait:** A process holding at least one resource is waiting to acquire additional resources held by other processes

Operating Systems 5.13 Part 5 Deadlock and Starvation

Deadlock Conditions (Cont.)

3. **No preemption:** A resource can be released only voluntarily by the process holding it, after that process has completed its task
4. **Circular wait:** There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for P_2, \dots, P_{n-1} is waiting for P_n , and P_n is waiting for P_0

Operating Systems 5.14 Part 5 Deadlock and Starvation

When there are multiple instances per resource type, we are not sure whether the graph has deadlock even if the graph has cycles. To handle this case, we study the deadlock conditions, that is, conditions necessary for a deadlock to occur. Deadlock can arise if the four conditions hold simultaneously. If any of them is violated, there is no deadlock. Also, the four conditions are necessary, but not sufficient.

We will briefly talk about the definitions here, and then give more examples in the later slides.

No preemption means we cannot take the resource from the process by force. Circular wait corresponds to a cycle in the resource allocation graph.

Part 5: Deadlocks and Starvation

- Deadlock Problem
- System Model
- Deadlock Conditions
- **Deadlock Prevention**
- Deadlock Avoidance
- Deadlock Detection

Operating Systems 5.15 Part 5 Deadlock and Starvation

Methods for Handling Deadlocks

- Ensure that the system will never enter a deadlock state
- Allow the system to enter a deadlock state and then recover
- Ignore the problem and pretend that deadlocks never occur in the system
 - Used by most operating systems, including UNIX

Operating Systems 5.16 Part 5 Deadlock and Starvation



There are multiple approaches to handling deadlocks.

First, we can ensure that the system will never enter a deadlock state. That is deadlock prevention and avoidance.

Second, we can allow the system to enter a deadlock state and then recover. That is deadlock detection.

Third, we simply ignore the problem and pretend that deadlocks never occur in the system. That means, do nothing. That is so called Ostrich approach. Why do most systems use Ostrich approach? Efficiency. The overhead of deadlock detection and prevention is high. Why do we learn about deadlock detection and prevention then? Because OS doesn't implement deadlock, it is the user's responsibility to handle deadlock in their programs. So, we need to pay attention to deadlock issues when we develop our multi-programmed software.

Deadlock Prevention

Prevent at least one of the four deadlock conditions
We will illustrate using Dining-Philosophers Problem

- Recall dining-philosopher solution using semaphores
 - Each chopstick is a shared resource protected by a binary semaphore (`chopstick [5]`);
 - Initially, for all i, `chopstick [i] = 1`;
 - Code (Process Philosopher i)



```

while (1) {
    wait(chopstick [ i ]);
    wait(chopstick [ (i+1)%5 ]);
    eat
    signal(chopstick [ i ]);
    signal(chopstick [ (i+1)%5 ]);
    think
}

```

5.17 Part 5 Deadlock and Starvation

Deadlock Prevention (Cont.)

Prevent at least one of the four deadlock conditions
We will illustrate using Dining-Philosophers Problem

- 1. Mutual Exclusion**
 - Chopsticks are not shareable (simultaneously), hence this condition cannot be eliminated
- 2. Hold and Wait:** Must guarantee that whenever a process requests a resource, it does not hold any other resource
 - Allow a philosopher to pick up chopsticks only if both the required chopsticks are available

need 1 semaphore
wait for mutex
me get or wait 1 chopstick, release mutex

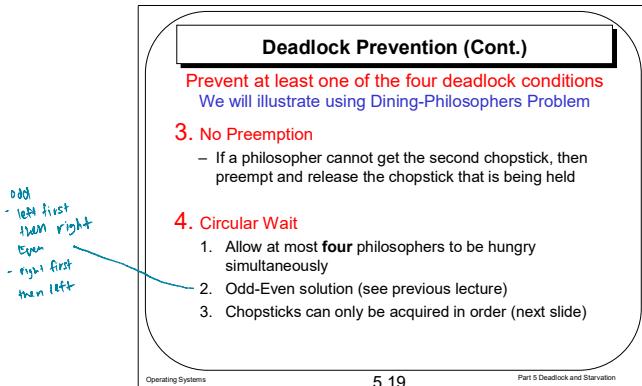
5.18 Part 5 Deadlock and Starvation

For philosopher i ($0 \leq i < 5$): He first picks up his left chopstick, and then the right chopstick. This is done by calling the wait operation on the corresponding semaphore. If both operations can succeed, the philosopher can eat. After eating, he releases the left chopstick and then the right one. After that, the philosopher will think, and then repeat the process.

Does this solution have any problem? Each philosopher executes the first line and a context switch occurs. → deadlock.

We will talk about how to prevent the deadlock conditions one by one. Recall, the four conditions are necessary to hold simultaneously for a deadlock to occur. So, we can prevent any one of them to happen. We illustrate that with Dining-Philosophers as an example.

Hold and wait: In Dining-Philosophers, allow a philosopher to pick up his chopsticks only if both the chopsticks are available. That means, at the beginning, the philosopher hasn't taken any chopstick. In that case, there is no hold and wait. So, no deadlock.



odd
- left first
even
- right first
then left

3. No Preemption

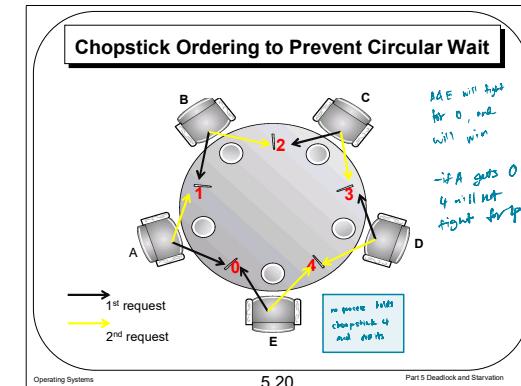
- If a philosopher cannot get the second chopstick, then preempt and release the chopstick that is being held

4. Circular Wait

1. Allow at most **four** philosophers to be hungry simultaneously
2. Odd-Even solution (see previous lecture)
3. Chopsticks can only be acquired in order (next slide)

case 1 :
 fight for 0 → A & E
 → suppose E wins, get 0
 → Next fight for 4
EITHER D or E will get to eat

case 2 :
 fight for 0 → A & E
 → Suppose A wins, get 0
 → A can eat it A gets 1
 → it not, B gets 1
 → B can eat when it get 2
 → it not, C can eat it get 3
 → it not, D guaranteed to eat
 ∴ A, B, C or D guaranteed to eat
 → cause E not fighting for 4



In the dining philosopher problem, each chopstick has an ID. Suppose each philosopher takes the chopstick with the small id first, and then takes the chopstick with the large id.

First request in black.

Second request in yellow.

→ No circular wait.

→ No deadlock.

Circular wait: In Dining-Philosophers, allow at most **four** philosophers to be hungry simultaneously. This is because of pigeon hole theory.

Part 5: Deadlocks and Starvation

- Deadlock Problem
- System Model
- Deadlock Conditions
- Deadlock Prevention
- **Deadlock Avoidance**
- Deadlock Detection

Operating Systems 5.21 Part 5 Deadlock and Starvation

Deadlock Avoidance

- The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that the system never goes into **unsafe state**
 - When a process requests an available resource, system must decide if immediate allocation leaves the system in a **safe state**
 - If **safe**, the request is granted, and otherwise, the process must wait
- System is in safe state if there exists a **safe completion sequence** of all processes without deadlock

Operating Systems 5.22 Part 5 Deadlock and Starvation

The deadlock prevention breaks one of the deadlock conditions. We now consider a more general deadlock avoidance algorithm. The algorithm is run whenever a process requests resources. The algorithm examines the resource-allocation state, and ensures that the system never goes into unsafe state. That is, when a process requests an available resource, system must decide if immediate allocation leaves the system in a **safe state**. If **safe**, the request is granted. Otherwise, the process must wait.

So what is safe state? System is in safe state if there exists a **safe sequence** of all processes. We do not require all sequences are safe. Only one safe sequence is sufficient.

Safe State

- A process completion sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe, if for each P_i , the resources that P_i requests can be satisfied by currently available resources + resources held by all the $P_j, j < i$
 - If P_i 's needs cannot be immediately met, then P_i can wait until all $P_j, j < i$ have finished
 - When all $P_j (j < i)$ are finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
 - When P_i terminates, P_{i+1} can obtain its resources, and so on
 - All processes in the sequence can finish

Operating Systems 5.23 Part 5 Deadlock and Starvation

Example of Safe Process Sequence

Processes	Hold	Request
P1	1	1
P2	1	2
P3	1	3

Available: 1

Q1: Is $\langle P1, P2, P3 \rangle$ safe?
Yes
P1.request<Available
P2.request<=P1.Hold+Available
P3.request<=P1.Hold+P2.Hold+Available

Q2: $\langle P3, P2, P1 \rangle$ safe?
No
P3.request>Available

Q3: Is the system safe?
Yes
Exists one safe sequence
 $\langle P1, P2, P3 \rangle$

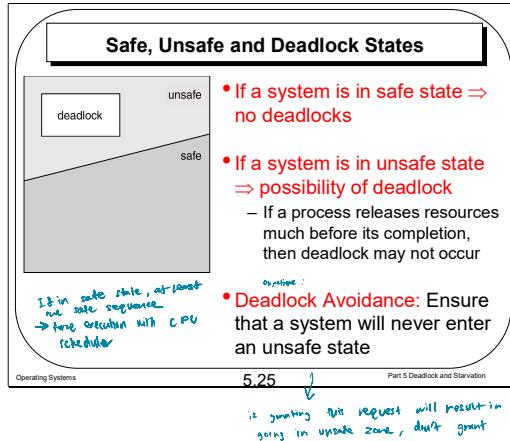
Operating Systems 5.24 Part 5 Deadlock and Starvation

Sequence $\langle P_1, P_2, \dots, P_n \rangle$ is safe if for each P_i , the resources that P_i requests can be satisfied by currently available resources + resources held by all the P_j with $j < i$. That is the processes before P_i in the sequence.

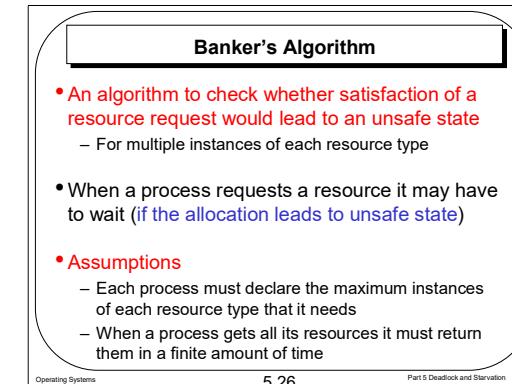
Consider a system with a single resource type. The number of instances available is one. There are three processes. P1 holds one instance and requests for one more and so on.

As an OS, what it knows is the resources they are holding and what they will need in future.
It doesn't know when they are going to request for resources in their execution, and when they are going to release resources in their execution.

Assumption made: resource held by process till the end of execution



The concept of bankers algo will be examinable on algo test!



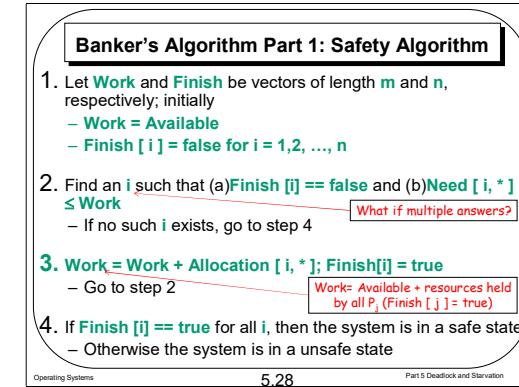
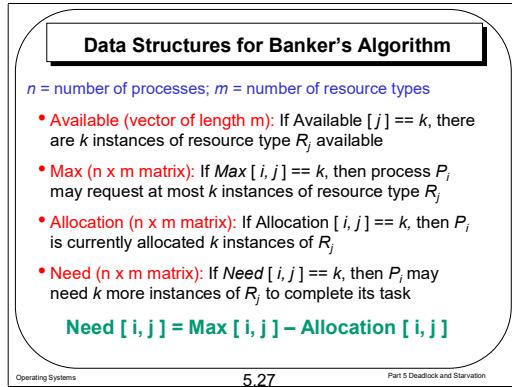
By definition, if a system is in a safe state, no deadlock.

If the system is in unsafe state, that means, there is no safe sequence for the system. It only means possibility of a deadlock. Because we are considering the worst case in the safe sequence: the resources are released only when the process finishes (**the worst case**). If the process can release resources before termination. The later processes in the sequence can actually release the resources earlier. So the previous process have more resources to consume. The safe sequence does not consider this case. So, the unsafe state does not mean deadlock. For example, you can change the example in the previous slide: available = 0. the system is unsafe. But, if P2 releases the resource earlier, then, all the processes can finish execution.

Now, let's learn the banker's algorithm for deadlock avoidance. The algorithm decides whether to meet a resource request, and to avoid the unsafe state. By its name: there are many bankers. They can borrow and lend money to each other. The banker's algorithm is to make sure that the bankers will not have a deadlock problem.

It can be used for multiple instances of each resource type. If all resource types have a single instance, what algorithm can we use? We can use resource-allocation graph, and see if the resource-allocation graph has a cycle.

We have some assumptions on the process execution. Before execution, each process must declare the maximum number of instances of each resource type that it needs. When a process requests a resource it may have to wait. that is when the allocation leads to unsafe state. When a process gets all its resource, it must return them in a finite amount of time. That means, eventually, the process will finish, and the waiting time of other processes is bounded.



Before we introduce the banker's algorithm, we define the data structures it uses.

Let's first introduce safety algorithm. The algorithm determines whether the current system is in a safe state or not. If the system is in a safe state, we will find a safe sequence for the system.

In the first step, we define two vectors work and finish to denote the current system state. Work represents the available resources during the execution of the algorithm. We will update it in Step 3. Finish denotes whether a process can be finished. If $\text{Finish}[i]$ is true, this means P_i can finish. It is also updated in Step 3.

What if multiple answers? Any one can be chosen. If we go to step 3 and check step 2 again, those processes still satisfy the condition.

Example of Safety Algorithm			
<ul style="list-style-type: none"> • 5 processes: P_0 through P_4 • 3 resource types: A (10 instances), B (5 instances) and C (7 instances) • Snapshot at time T_0: <u>Allocation</u> <u>Max</u> <u>Available</u> 			
	A B C	A B C	A B C
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

Operating Systems

5.29

Part 5 Deadlock and Starvation

Example of Safety Algorithm (Cont.)			
<ul style="list-style-type: none"> • Then the matrix Need is defined to be Max – Allocation 			
	<u>Allocation</u>	<u>Max</u>	<u>Need</u>
	A B C	A B C	A B C
P_0	0 1 0	7 5 3	Max - Allocation \rightarrow 7 4 3
P_1	2 0 0	3 2 2	1 2 2
P_2	3 0 2	9 0 2	6 0 0
P_3	2 1 1	2 2 2	0 1 1
P_4	0 0 2	4 3 3	4 3 1

Operating Systems

5.30

Part 5 Deadlock and Starvation

Now, let's look at one example. There are five processes, P_0 to P_4 . There are 3 resource types: A with 10 instances, B with 5 instances, and C with 7 instances.

The system state at time T_0 is as follows.

Allocation

Max

Available

For example: P_0 is allocated with 1 instance of B, and P_0 may request 7 instances of A, and 5 instances of B, and 3 instances of C at most.

So, now, let's calculate Need. By definition, $\text{Need} = \text{Max} - \text{Allocation}$.

possible to have
many complex sequences

Example of Safety Algorithm (Cont.)

- The working is as follows:

<u>Need</u>	<u>Allocation</u>	Work			Finished Process
A B C	A B C	A	B	C	
P ₀ 7 4 3 ✓	0 1 0	+	2	3	P ₁
P ₁ 1 2 2 ✓	2 0 0		5	3	P ₂
P ₂ 6 0 0 ✓	3 0 2	+	2	1	P ₃
P ₃ 0 1 1 ✓	2 1 1	+	0	0	P ₄
P ₄ 4 3 1 ✓	0 0 2	+	3	7	P ₂
The system is safe (sequence <P₁, P₃, P₄, P₂, P₀>)		10	4	7	
		+	0	1	P ₀
			10	5	7

doesn't matter
which one
we start with — ①

Banker's Part 2: Resource-Request Algorithm

Request_i = Request vector for process **P_i**
 – If **Request_i[j] == k** then process **P_i** wants **k** instances
 of resource type **R_j**

$V1 \leftarrow V2 \equiv V1[j] \leftarrow V2[j], \text{ for all } j$

→ The i^{th} row of the matrix

1. If $\text{Request}_i \leq \text{Need}_i$, go to step 2
 - Otherwise, raise error condition, since process has exceeded its maximum claim
 2. If $\text{Request}_i \leq \text{Available}$, go to step 3
 - Otherwise P_i must wait, since resources are not available

Is there any other completion sequence? As long as you can find a completion sequence for the process, we can say the system is safe.

We now describe the algorithm for resource-request algorithm for a process P_i . We decide whether to grant the resource to the process or let the process wait.

The resource request by the process P_i is represented in a vector Request . If $\text{Request}[j] = k$ then process P_i wants k instances of resource type R_j .

Resource-Request Algorithm (Cont.)

3. Pretend to allocate requested resources to P_i by modifying the state as follows:
 - Available = Available – Request_i;
 - Allocation_i = Allocation_i + Request_i;
 - Need_i = Need_i – Request_i;
4. Run the safety algorithm
 - Safe \Rightarrow the resources are allocated to P_i
 - Unsafe $\Rightarrow P_i$ must wait, and the old resource-allocation state is restored

Operating Systems 5.33 Part 5 Deadlock and Starvation

Example of Resource-Request Algorithm

- Suppose P_1 requests (1, 0, 2) resources
- Step 1: Check Request₁ \leq Need₁, (1,0,2) \leq (1,2,2) \Rightarrow true
- Step 2: Check that Request₁ \leq Available, i.e. (1,0,2) \leq (3,3,2) \Rightarrow true
- Steps 3 & 4:

	Allocation	Max	Need	Available
A	0 1 0	7 5 3	7 4 3	3 3 2
B	2 0 0	3 2 2	1 2 2	2 3 0
C	3 0 2	9 0 2	0 2 0	6 0 0
P_0				
P_1				
P_2				
P_3				
P_4				

Step 3: Pretend to allocate requested resources
 Step 4: Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ is safe

Operating Systems 5.34 Part 5 Deadlock and Starvation

In step 3, we pretend to allocate the requested resources.
 What are the parameters that we need to change?
 Do we need to change Max? no.

Example of Resource-Request Algorithm

- Can a further request for $(1,0,2)$ by P_1 be granted?
 - Step 1: Check if $\text{Request}_1 \leq \text{Need}_1$, $(1,0,2) \not\leq (0,2,0) \Rightarrow \text{false}$
- Can a further request for $(0,2,0)$ by P_0 be granted?
 - Step 1: Check if $\text{Request}_0 \leq \text{Need}_0$, $(0,2,0) \leq (7,4,3) \Rightarrow \text{true}$
 - Step 2: Check if $\text{Request}_0 \leq \text{Available}$, $(0,2,0) \leq (2,3,0) \Rightarrow \text{true}$

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	ABC	ABC	ABC
P_0	0 1 0	7 4 3	2 3 0
P_1	0 3 0	7 2 3	0 2 0
P_2	3 0 2	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

• How about a request for $(2,3,0)$ by P_4 ?

@We cannot find $\text{Need}_4 < \text{Available}$
→ Restore the state

Operating Systems 5.35 Part 5 Deadlock and Starvation

it allows,
it will
violate
safe state

NOT EXAMINABLE

Part 5: Deadlocks and Starvation

- Deadlock Problem
- System Model
- Deadlock Conditions
- Deadlock Prevention
- Deadlock Avoidance
- **Deadlock Detection**

Operating Systems 5.36 Part 5 Deadlock and Starvation

Deadlock Detection

- Allow system to enter deadlock state
- Then invoke detection algorithms
 - For single instance of each resource type (in textbook based on identifying cycle in resource-allocation graph)
 - For multiple instances of each resource type (next few slides)
- Then invoke recovery algorithm (in textbook)

This slide and the later slides in this chapter are not examinable

Operating Systems 5.37 Part 5 Deadlock and Starvation

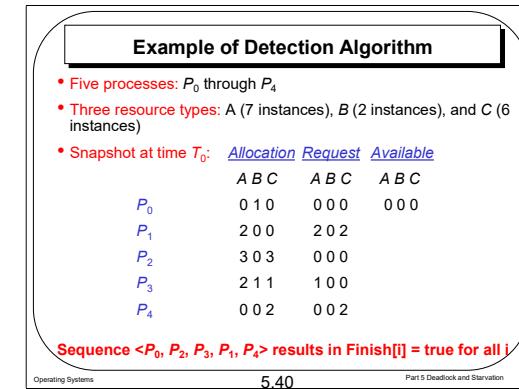
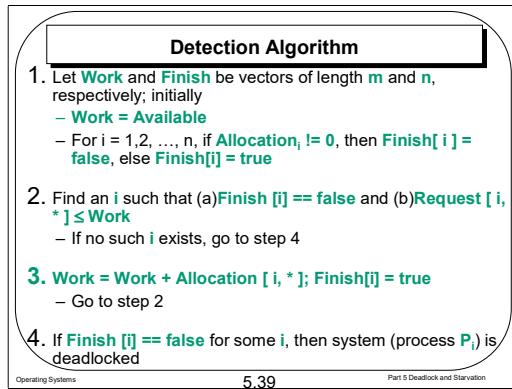
Multiple Instances of Each Resource Type

- **Available:** A vector of length m indicates the number of available resource types
- **Allocation:** An $n \times m$ matrix defines the number of resources of each type currently allocated to each process
- **Request:** An $n \times m$ matrix indicates the current request of each process. If $\text{Request}[i, j] = k$, then process P_i is requesting k more instances of resource type R_j

Operating Systems 5.38 Part 5 Deadlock and Starvation

The basic idea of deadlock detection is to allow system to enter deadlock state, then invoke detection algorithm periodically. If there is a deadlock, we invoke the recovery algorithm.

The deadlock detection algorithm is actually similar to deadlock avoidance algorithm. If you are interested in them, you can refer to textbooks or just go through them in slides. They are not examinable.



Let's first introduce safety algorithm. The algorithm determines whether the current system is in a safe state or not. If the system is in a safe state, we will find a safe sequence for the system.

In the first step, we define two vectors work and finish to denote the current system state. Work represents the available resources during the execution of the algorithm. We will update it in Step 3. Finish denotes whether a process can be finished. If **Finish[i]** is true, this means P_i can finish. It is also updated in Step 3.

What if multiple answers? Any one can be chosen. If we go to step 3 and check step 2 again, those processes still satisfy the condition.

Example of Detection Algorithm (Cont.)																							
Suppose P_2 requests an additional instance of type C																							
<table style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th style="text-align: left; width: 33.33%;"><u>Allocation</u></th> <th style="text-align: left; width: 33.33%;"><u>Request Available</u></th> <th style="text-align: left; width: 33.33%; text-align: center;">Work Finished</th> </tr> </thead> <tbody> <tr> <td style="text-align: left; vertical-align: bottom;">$A \ B \ C$</td> <td style="text-align: left; vertical-align: bottom;">$A \ B \ C$</td> <td style="text-align: left; vertical-align: bottom;">$A \ B \ C$</td> </tr> <tr> <td style="text-align: left; vertical-align: bottom;">$P_0 \ 1 \ 0$</td> <td style="text-align: left; vertical-align: bottom;">$0 \ 0 \ 0$ ✓</td> <td style="text-align: left; vertical-align: bottom;">$0 \ 0 \ 0$</td> </tr> <tr> <td style="text-align: left; vertical-align: bottom;">$P_1 \ 2 \ 0 \ 0$</td> <td style="text-align: left; vertical-align: bottom;">$2 \ 0 \ 2$ ⊗</td> <td style="text-align: left; vertical-align: bottom;">$0 \ 1 \ 0$</td> </tr> <tr> <td style="text-align: left; vertical-align: bottom;">$P_2 \ 3 \ 0 \ 3$</td> <td style="text-align: left; vertical-align: bottom;">$0 \ 0 \ 1$ ⊗</td> <td style="text-align: left; vertical-align: bottom;"></td> </tr> <tr> <td style="text-align: left; vertical-align: bottom;">$P_3 \ 2 \ 1 \ 1$</td> <td style="text-align: left; vertical-align: bottom;">$1 \ 0 \ 0$ ⊗</td> <td style="text-align: left; vertical-align: bottom;"></td> </tr> <tr> <td style="text-align: left; vertical-align: bottom;">$P_4 \ 0 \ 0 \ 2$</td> <td style="text-align: left; vertical-align: bottom;">$0 \ 0 \ 2$ ⊗</td> <td style="text-align: left; vertical-align: bottom;"></td> </tr> </tbody> </table>			<u>Allocation</u>	<u>Request Available</u>	Work Finished	$A \ B \ C$	$A \ B \ C$	$A \ B \ C$	$P_0 \ 1 \ 0$	$0 \ 0 \ 0$ ✓	$0 \ 0 \ 0$	$P_1 \ 2 \ 0 \ 0$	$2 \ 0 \ 2$ ⊗	$0 \ 1 \ 0$	$P_2 \ 3 \ 0 \ 3$	$0 \ 0 \ 1$ ⊗		$P_3 \ 2 \ 1 \ 1$	$1 \ 0 \ 0$ ⊗		$P_4 \ 0 \ 0 \ 2$	$0 \ 0 \ 2$ ⊗	
<u>Allocation</u>	<u>Request Available</u>	Work Finished																					
$A \ B \ C$	$A \ B \ C$	$A \ B \ C$																					
$P_0 \ 1 \ 0$	$0 \ 0 \ 0$ ✓	$0 \ 0 \ 0$																					
$P_1 \ 2 \ 0 \ 0$	$2 \ 0 \ 2$ ⊗	$0 \ 1 \ 0$																					
$P_2 \ 3 \ 0 \ 3$	$0 \ 0 \ 1$ ⊗																						
$P_3 \ 2 \ 1 \ 1$	$1 \ 0 \ 0$ ⊗																						
$P_4 \ 0 \ 0 \ 2$	$0 \ 0 \ 2$ ⊗																						
Operating Systems	5.41	Part 5 Deadlock and Starvation																					

Example of Detection Algorithm (Cont.)		
<ul style="list-style-type: none"> • State of system? <ul style="list-style-type: none"> – Can reclaim resources held by process P_0, but insufficient resources to fulfill other process requests – Deadlock exists, consisting of processes P_1, P_2, P_3, and P_4 		
<p>Operating Systems</p>		
5.42		Part 5 Deadlock and Starvation