

Effective memory management is vital in a multiprogramming system. To improve CPU utilization, multiple processes need to be loaded into memory to ensure a reasonable supply of ready processes to consume available CPU time.

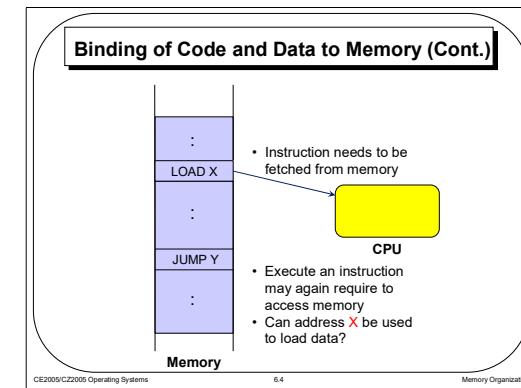
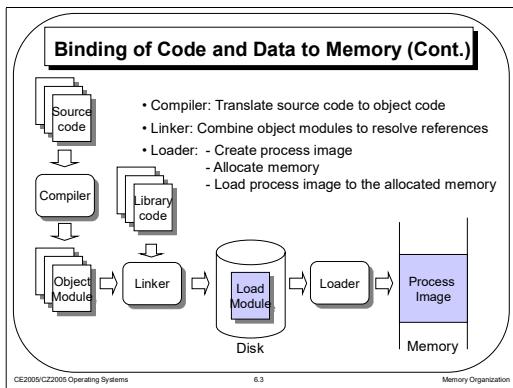
We begin this chapter by understanding how a process gets executed. Next, we approach the technology of memory management by looking at a variety of schemes that have been used.

How is memory allocated for a process image?

Logical view of a process in memory (also called process image) – not just code and data, but also include dynamic memory space (stack and heap) required during process execution

It defines logically the space a process is able to access (or address) during its execution. So, this is called logical address space of a process. Address reference generated when executing an instruction always refers to this logical address space and is called logical address.

*more definitions*



How is a process image generated and loaded?

Loader is part of OS

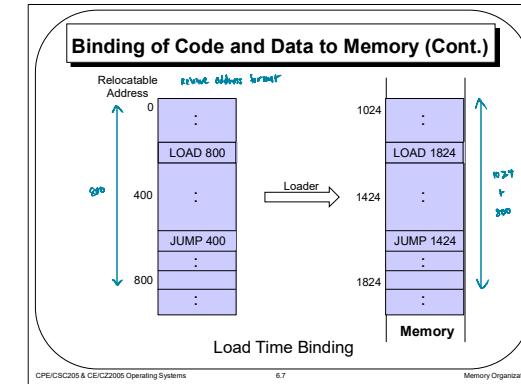
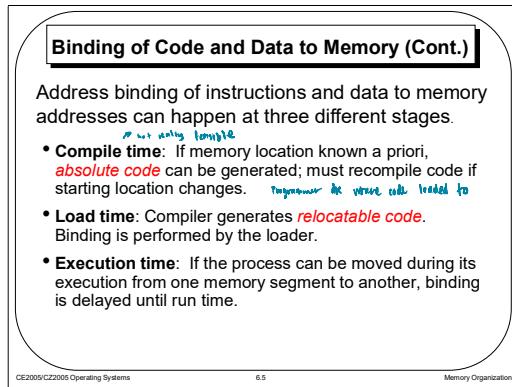
A program may contain multiple source code files. Source code files are compiled into object files. The linker combines these object files into a single binary executable (i.e., the load module on disk). During the linking phase, other object files or libraries may be included as well, such as the standard C or math library. A loader is used to load the binary executable file into memory where its is eligible to run on a CPU core.

Instruction execution cycle

As a process executes, it accesses instructions and data in memory

CPU must deal with memory references ~~within the program~~. Branch instructions contain an address to reference the instruction to be executed next. Data reference instructions contain the address of the byte or word of data referenced.

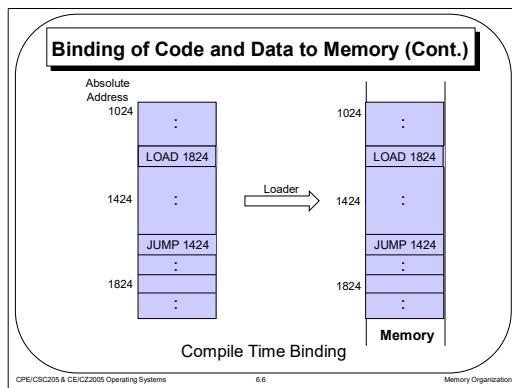
Address generated while executing an instruction may not be used directly to access physical memory, depending on when address binding is performed.

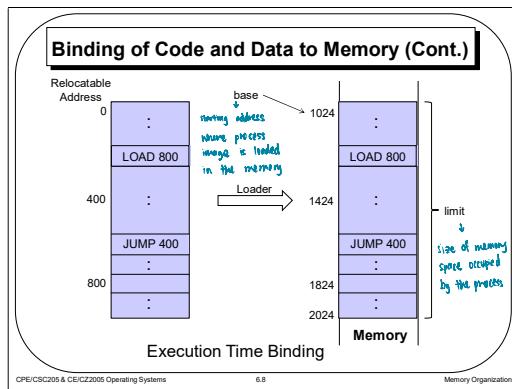


Address binding refers to the action of binding (or translating) the addresses in process image to the physical memory addresses (or absolute addresses).

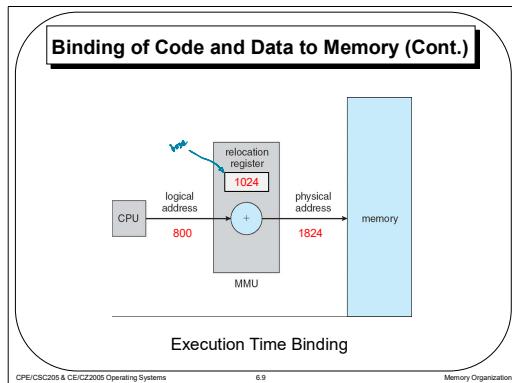
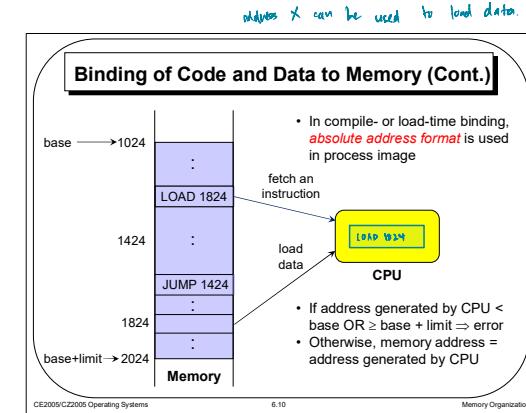
Compiler generates code in re-locatable address format, i.e., address format that is relative to the beginning of the address space (also called relative address).

Loader will translate the re-locatable address to absolute address.





Base – starting address where process image is loaded in the memory.  
Limit – size of the memory space occupied by the process.



For compile-time and load-time binding, absolute address is used directly to access memory.  
So, the binding of logical address space to the physical memory is static. Once process image is loaded, it can't be moved in the memory.

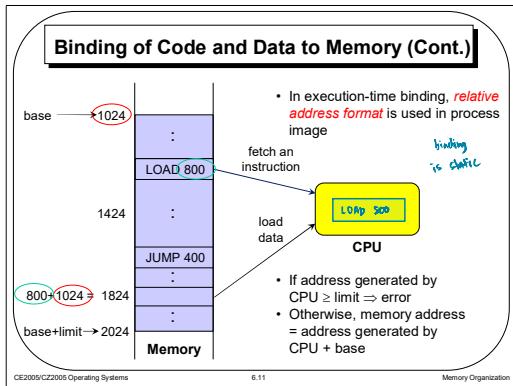
In a multiprogramming system, the available memory is generally shared among a number of processes. Typically, it is not possible for the programmer to know in advance which other programs will be resident in physical memory at the time of execution of his or her program. In addition, we would like to be able to:

- Perform memory compaction
- Swap in and out process
- Provide virtual memory support

Hence, we are not interested in compile- or load-time binding.

*address × cannot be used to load data.*

*Focus on  
execution  
time binding  
for this  
slide.*



### Logical vs. Physical Address Space

- **Address Space** – all addresses accessible by a process
- **Logical address** – address used in the code, generated by the CPU when executing an instruction.
- **Physical address** – address used to access physical memory, seen by the memory unit.

CE2009/CZ2005 Operating Systems 6.12 Memory Organization

Relative address – address format that is relative to the beginning of the process image.

Base – starting address where process image is loaded in the memory.

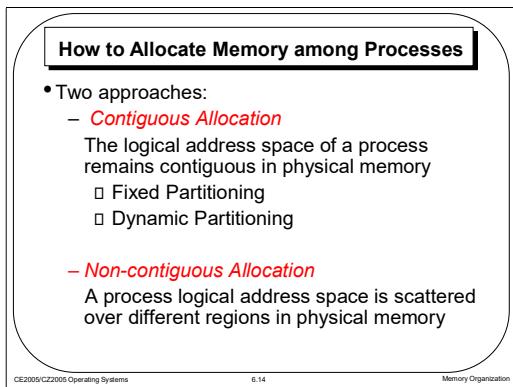
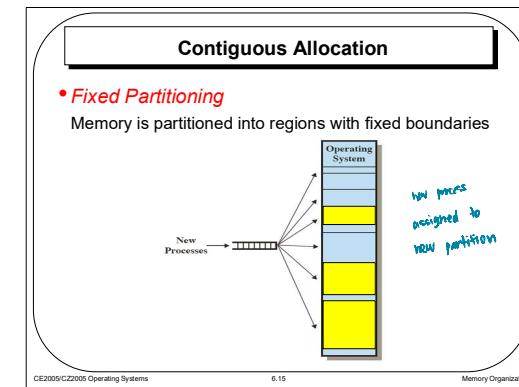
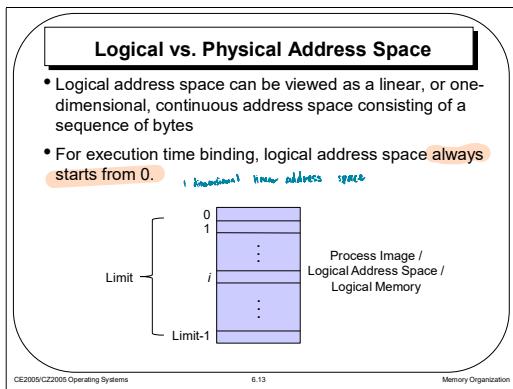
In execution-time binding, address generated by CPU needs to be translated to the one used to access memory and this translation depends on memory allocation method used.

The address translation is done by MMU and supported by hardware – recall base and limit registers.

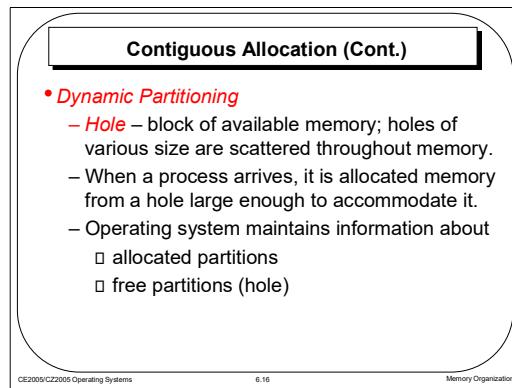
Execution-time binding is dynamic, process image can be moved in the memory. In the following lecture, we assume execution-time binding.

Process executes in logical address space and it never sees physical address. There is one logical address space per process.

To access physical memory, logical address may need to be translated into physical address.



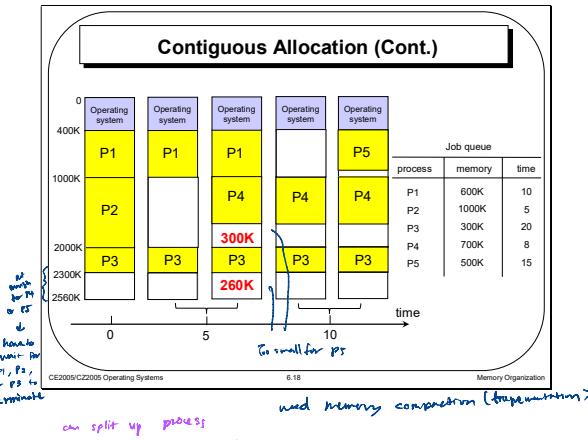
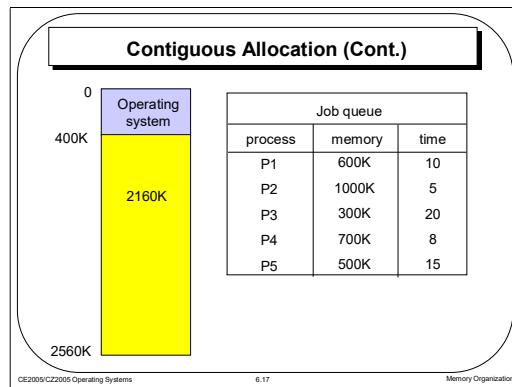
- The number of partitions specified at system generation time limits the number of active processes in the system.
- Because partition sizes are preset at system generation time, small processes will not utilize partition space efficiently. Any process, no matter how small, occupies an entire partition. This approach is not really practical. It is difficult to determine the number of partitions and size of each partition since the memory requirements of all processes are in general unknown beforehand.



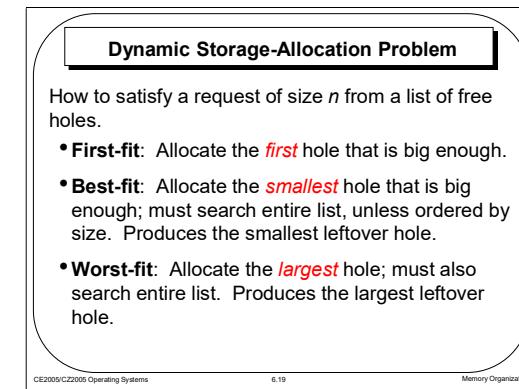
Initially, there is a single hole.

Memory allocation status may change while processes are created (allocate memory) and terminate (de-allocate memory).

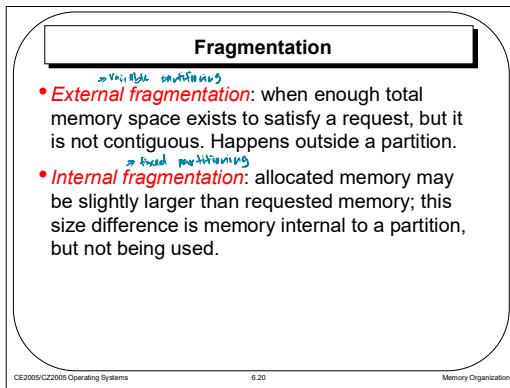
Eventually, as time goes, memory contains a list of holes of vary size.



With dynamic partitioning, the partitions are of variable length and number. When a process is brought into physical memory, it is allocated exactly as much memory as it requires and no more.

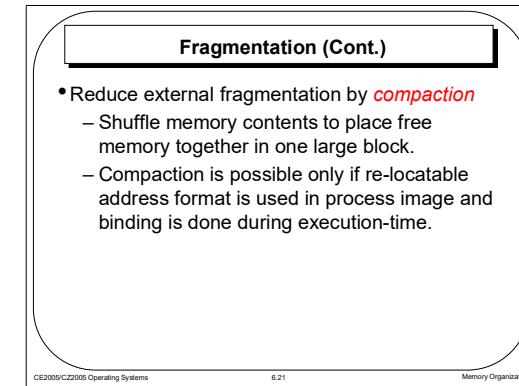


Given memory requirement of a process, dynamic storage allocation policies determine which one to use amongst a list of free holes.

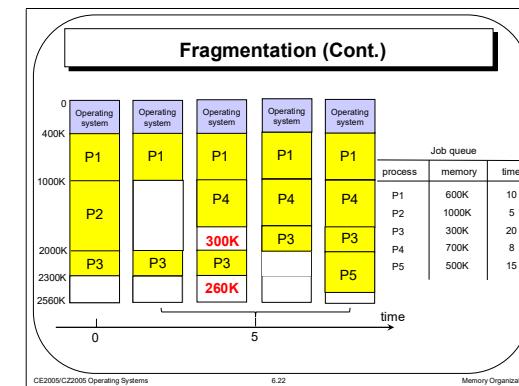


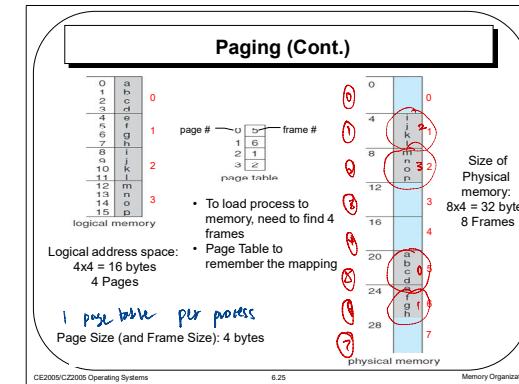
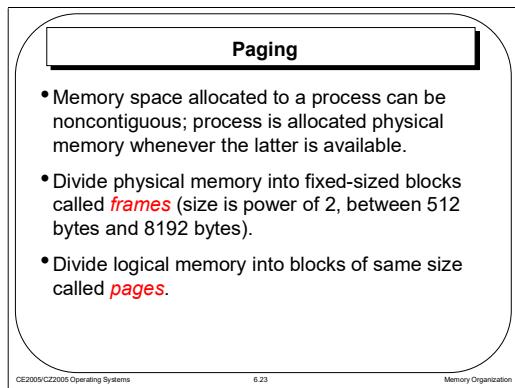
Variable partitioning -> external fragmentation

Fixed partitioning -> internal fragmentation



From time to time, the OS shifts the processes so that all of the free memory is together in one block. The difficulty with compaction is that it is a time-consuming procedure and wasteful of processor time. Note that compaction implies the need for a dynamic relocation capability. That is, it must be possible to move a process from one region to another in memory without invalidating the memory references in the program.





Memory space allocated to a process is not contiguous.

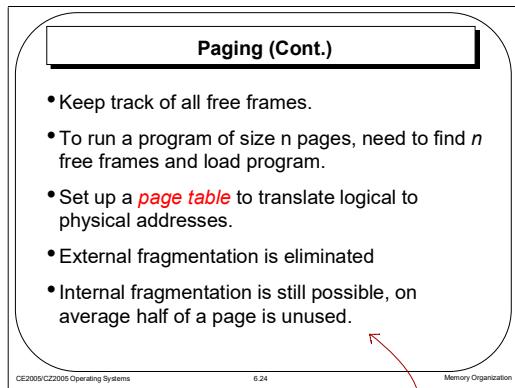
This is how paging works.

Example: LOAD 0111 (7)  
load "h" from logical address 0111

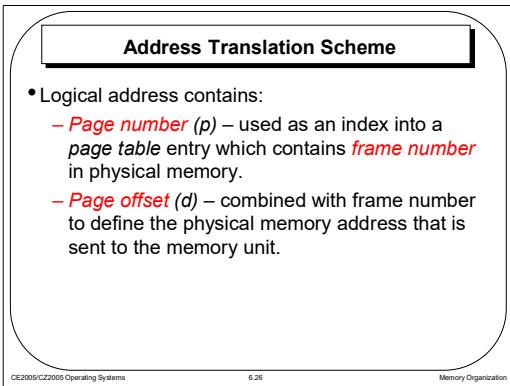
"h" is stored in physical location 11011 (27)

Translate "0111" to "11011"

To do address translation, need to know the logical address is in which page and the offset within the page

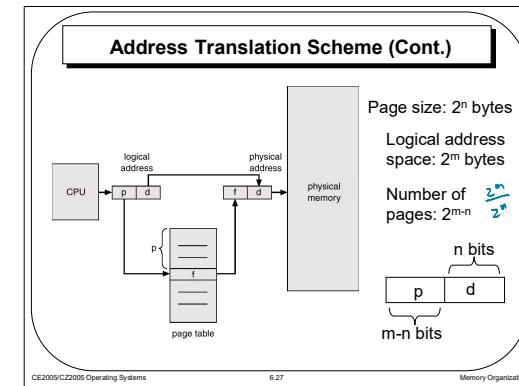


The operating system maintains a **page table** for each process. The page table shows the frame location for each page of the process.



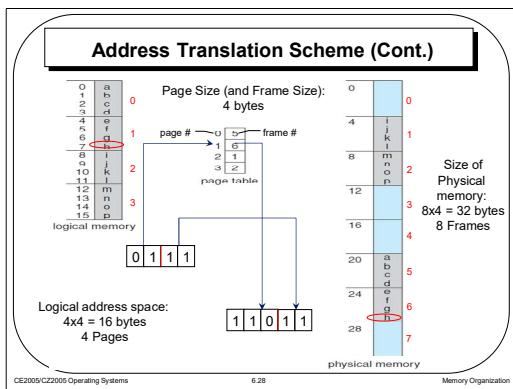
1 page 1 process  
1 TLB for all processes

Page size: 1024 bytes  $\rightarrow 2^{10}$  bytes  
Size of page table entry: 4 bytes  
logical address space:  $2^{32}$  bytes  
logical address space = no. of pages  
Page size  
no. of pages  $\times$  page table entry size  
 $= \frac{2^{32}}{2^10} \times 4 = 2^{20} \times 2^2 = 2^{22} = 4 \text{ MB}$



Logical address:  $m$  bits. There are  $2^m$  bytes in logical address space.  
Page size:  $2^n$ . There are  $2^m/2^n$  pages in logical address space.  
(For the convenience of addressing, page size is always power of 2.)  
Logical address is divided into two parts: lower  $n$  bits for page number and higher  $(m-n)$  bits for page number.

This figure suggests a hardware implementation. When a particular process is running, a register holds the starting address of the page table for that process. The page number of a virtual address is used to index that table and look up the corresponding frame number. This is combined with the offset portion of the virtual address to produce the desired real address. Typically, the page number field is longer than the frame number field.

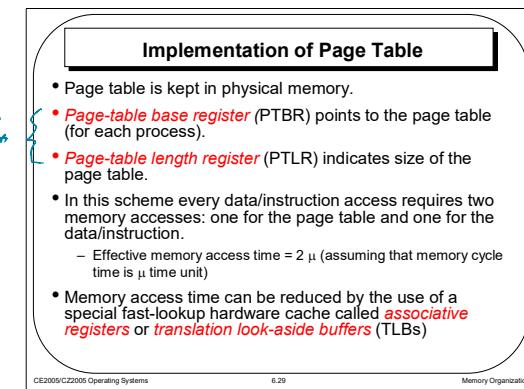


Example: LOAD 0111 (7)  
load "h" from logical address 0111

"h" is stored in physical location 11011 (27)

Translate "0111" to "11011"

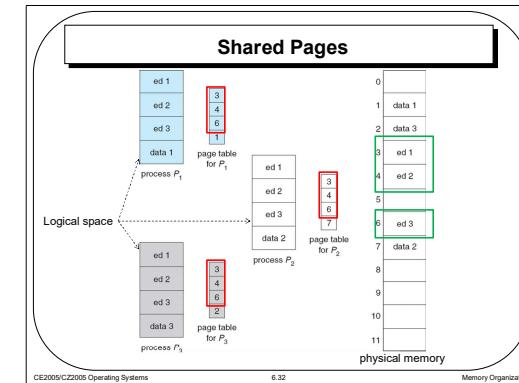
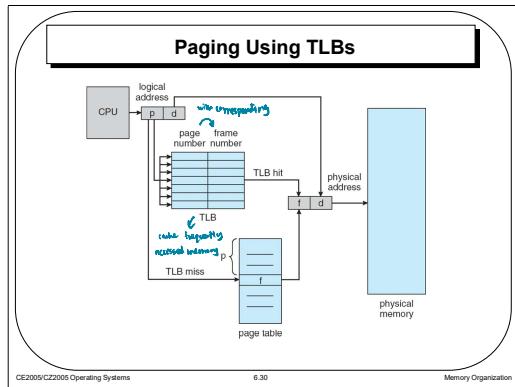
To do address translation, need to know the logical address is in which page and the offset within the page



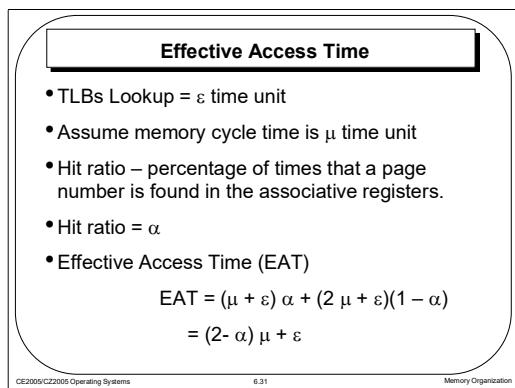
PAGE TABLE STRUCTURE The basic mechanism for reading a word from memory involves the translation of a virtual, or logical, address, consisting of page number and offset, into a physical address, consisting of frame number and offset, using a page table. Because the page table is of variable length, depending on the size of the process, we cannot expect to hold it in registers. Instead, it must be in physical memory to be accessed.

With paging, the logical-to-physical address translation is still done with hardware support. Now the CPU must know how to access the page table of the current process -> page-table base register and page-table length register.

In principle, every virtual memory reference can cause two physical memory accesses: one to fetch the appropriate page table entry and one to fetch the desired data. Thus, a straightforward virtual memory scheme would have the effect of doubling the memory access time. To overcome this problem, most virtual memory schemes make use of a special high-speed cache for page table entries, usually called a translation lookaside buffer (TLB).



Essentially, TLB are associate registers which support parallel search.  
Typically, there are 64 to 1024 entries in TLB.



Sharing pages will reduce the memory usage.

In this example, three processes execute the same editor but on different data. Each process has its own logical address space, but they contain the same editor code (i.e., pages ed1, ed2, and ed3). So, instead of loading three copies of the editor code, only one copy needs to be loaded in the memory and this copy can be shared amongst 3 processes (by mapping ed1, ed2, and ed3 in each process's logical address space to the corresponding editor pages loaded in the memory as shown in the slide). This is only possible if code is read-only code (re-entrant code).

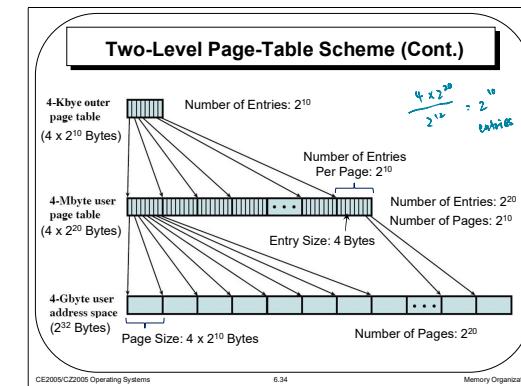
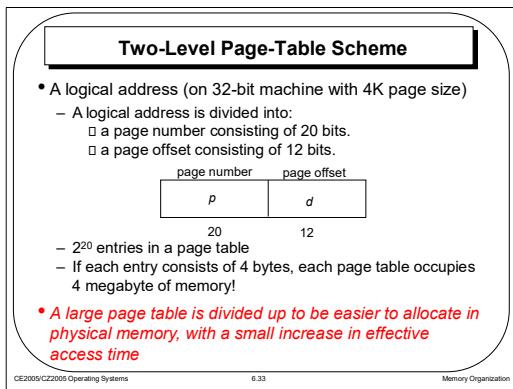
#### Shared code

One copy of read-only (reentrant) code shared among processes (i.e., text editors, compilers, window systems).

Shared code must appear in same location in the logical address space of all processes

#### Private code and data

Each process keeps a separate copy of the private code and data  
The pages for the private code and data can appear anywhere in the logical address space

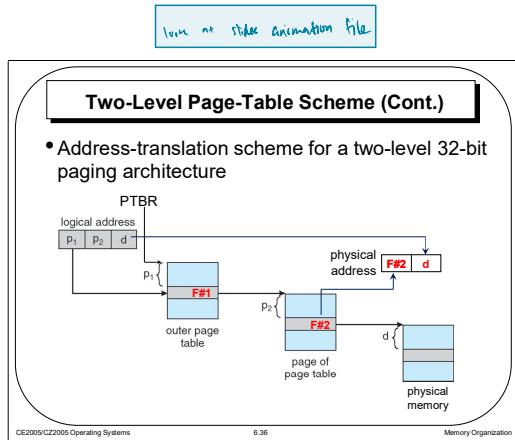
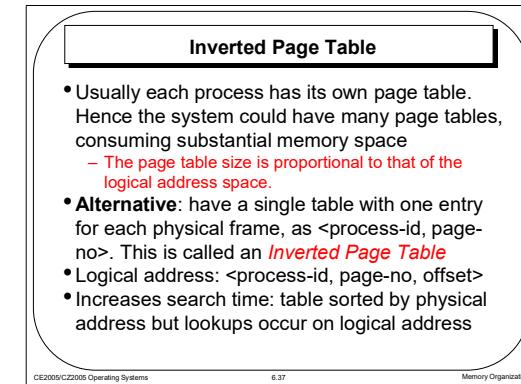
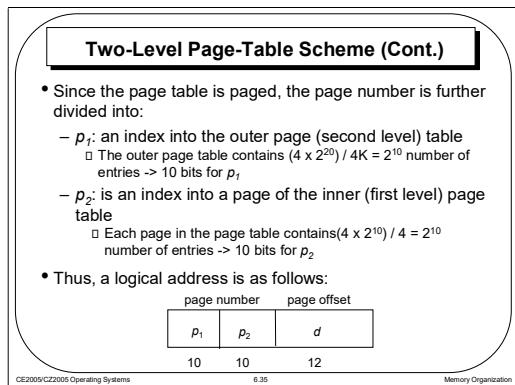


This figure shows an example of a two-level scheme typical for use with a 32-bit address.

If we assume byte-level addressing and 4-kbyte ( $2^{12}$ ) pages, then the 4-Gbyte ( $2^{32}$ ) virtual address space is composed of  $2^{20}$  pages.

If each of these pages is mapped by a 4-byte page table entry, we can create a user page table composed of  $2^{20}$  entries, requiring 4 Mbytes ( $2^{22}$ ).

This huge user page table, occupying  $2^{10}$  pages, can be kept in virtual memory and mapped by an outer page table with  $2^{10}$  entries, occupying 4 Kbyte ( $2^{12}$ ) of physical memory.



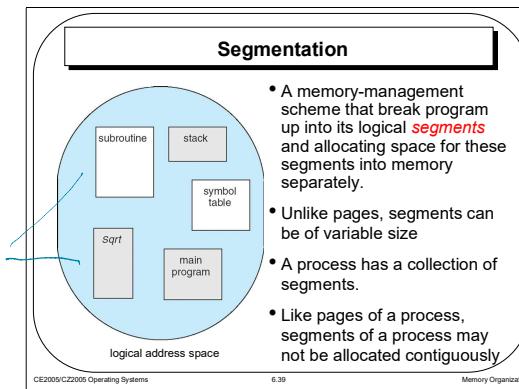
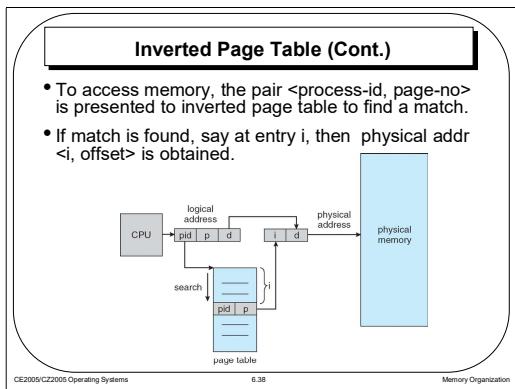
This figure shows the steps involved in address translation for this scheme. The outer page table always remains in physical memory. The first 10 bits of a virtual address are used to index into the outer page table to find an entry for a page of the page table. If that page is not in physical memory, a page fault occurs. If that page is in physical memory, then the next 10 bits of the virtual address index into the page table to find the entry for the page that is referenced by the virtual address.

Multiple-level paging is used in many systems, e.g., Intel IA-32 (Pentium) and IA-64 (Itanium).

A drawback of the type of page tables that we have been discussing is that their size is proportional to that of the logical address space and there is one page-table for each process.

An alternative approach is the use of an inverted page table structure.

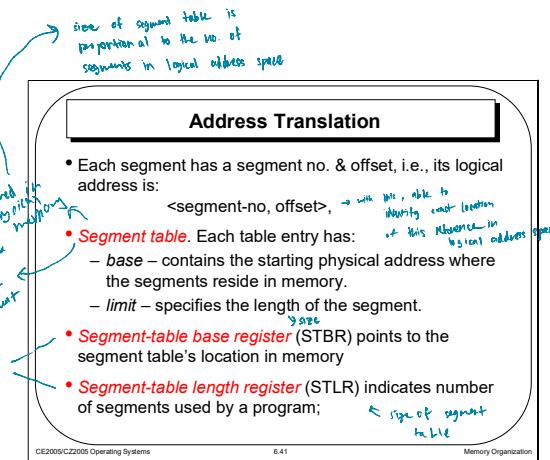
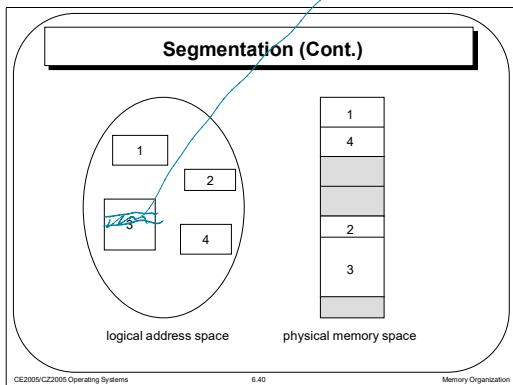
Entries in the inverted page table are ordered by the frame number order. Each entry contains <pid, page number> tuple. pid is needed because the inverted page table is shared by all the processes – the OS needs to know the page mapped to the frame belongs to which process.



Process id is added to logical address. Given the pair <pid, p>, the inverted page table is searched to locate a matching entry. If entry i is a matching entry (that is, it contains <pid, p>), the page p of process pid is mapped to frame i.

There is overhead in locating an entry, since entries are arranged in the frame number order (rather in the page number order as in conventional page table).

Logical address space is viewed as a collection of segments. A segment is a logical unit such as:  
main program, function, local variables, global variables, stack, symbol table, arrays



This slide shows logical address space of a process with 4 segments. The 4 segments are loaded to different places in the physical memory.

Segmentation has a number of advantages to the paging scheme:

**• It simplifies the handling of growing data structures.**

- With segmented virtual memory, the dynamic data structure (e.g., heap) can be assigned its own segment, and the operating system will expand or shrink the segment as needed.
- If a segment that needs to be expanded is in physical memory and there is insufficient room, the operating system may move the segment to a larger area of physical memory, if available, or swap it out. The enlarged segment would be swapped back in at the next opportunity.

**[2. It allows programs to be altered and recompiled independently,**

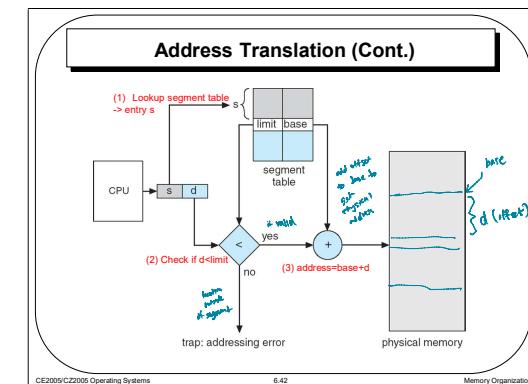
- without requiring the entire program to be relinked and reloaded.]

**3. It lends itself to sharing among processes.**

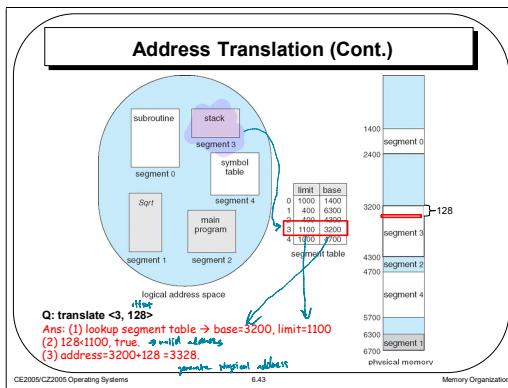
- Shared code/data can be put in a segment.

**[4. It lends itself to protection.**

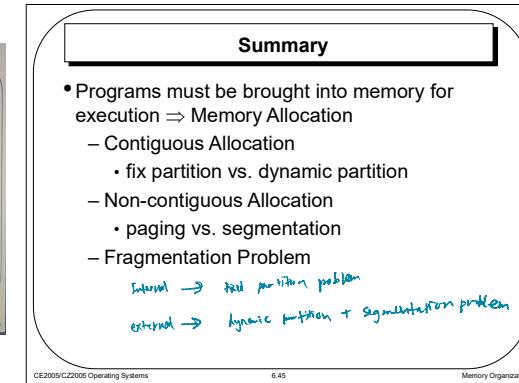
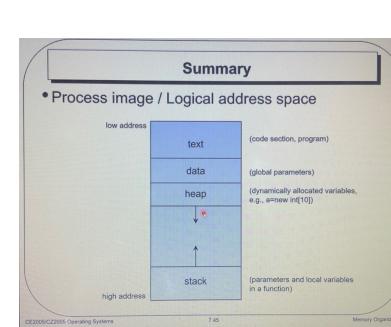
- Access right can be easily defined for each segment.]



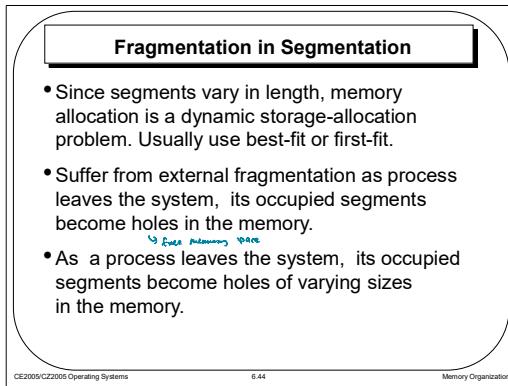
Slide 43



Slide 45

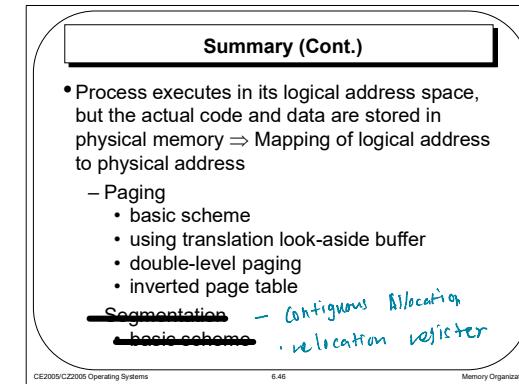


Slide 44

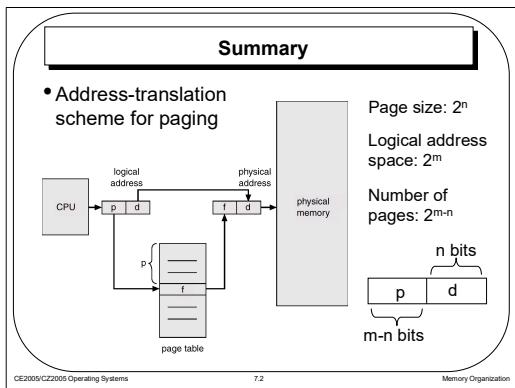
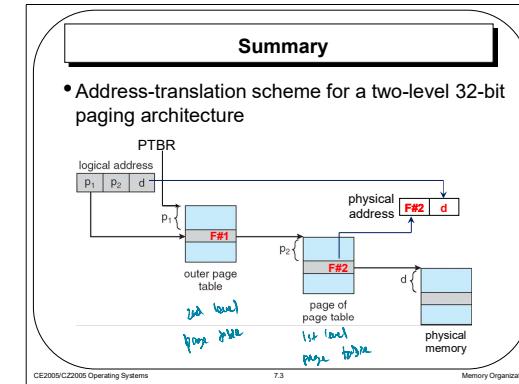
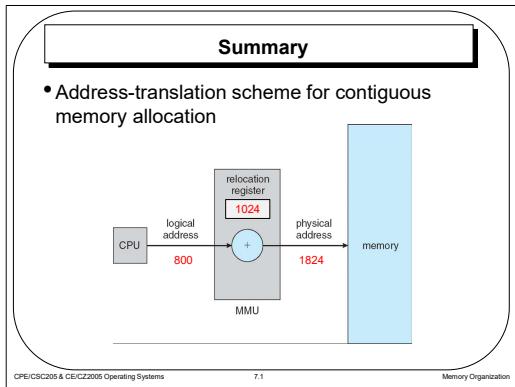


Memory allocation of a segmentation is similar to dynamic partitioning. Various dynamic storage allocation strategies (e.g., first-fit, best-fit, and worst-fit) can be used. The difference, compared to dynamic partitioning, is that a process image may occupy more than one partition, and these partitions need not be contiguous. Segmentation eliminates internal fragmentation but, like dynamic partitioning, it suffers from external fragmentation. However, because a process is broken up into a number of smaller pieces (segments), the external fragmentation should be less.

Slide 46

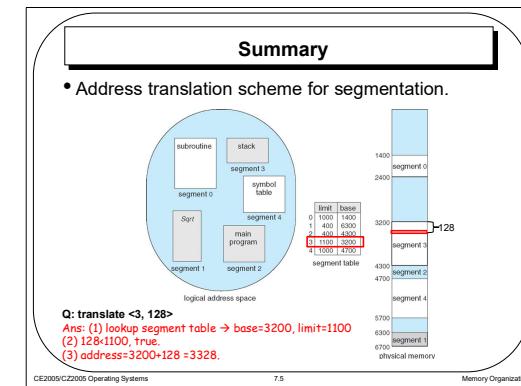
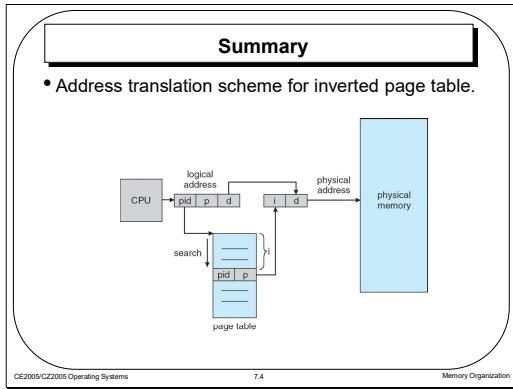


fixed or dynamic  
partitioning



This figure shows the steps involved in address translation for this scheme. The outer page table always remains in main memory. The first 10 bits of a virtual address are used to index into the outer page table to find an entry for a page of the page table. If that page is not in main memory, a page fault occurs. If that page is in main memory, then the next 10 bits of the virtual address index into the page table to find the entry for the page that is referenced by the virtual address.

This figure suggests a hardware implementation. When a particular process is running, a register holds the starting address of the page table for that process. The page number of a virtual address is used to index that table and look up the corresponding frame number. This is combined with the offset portion of the virtual address to produce the desired real address. Typically, the page number field is longer than the frame number field.



Process id is added to logical address. Given the pair  $\langle \text{pid}, p \rangle$ , the inverted page table is searched to locate a matching entry. If entry  $i$  is a matching entry (that is, it contains  $\langle \text{pid}, p \rangle$ ), the page  $p$  of process  $\text{pid}$  is mapped to frame  $i$ .

There is overhead in locating an entry, since entries are arranged in the frame number order (rather than in the page number order as in conventional page table).